

Supplementary Material

Multi-objective Pareto Hyperparameter Optimization for Joint Detection and Localization of Forged Identity Documents

Hernán Díaz Rodríguez

Department of Computer Science, University of Oviedo, Asturias, Spain

diazhernan@uniovi.es

This document is a companion to the Pattern Recognition Letters submission *Multi-objective Pareto Hyperparameter Optimization for Joint Detection and Localization of Forged Identity Documents* (DocVerify). It collects the complete experimental tables that, owing to the journal’s strict page limit, are summarized in only one or two sentences in the main paper or in the response to reviewers. All numbers are computed from the released result CSVs (archived together with this document); section references (e.g. Sec. 4.3) point to the main paper. Unless stated otherwise, all blind-test results use the held-out 15% FantasyID partition, 224×224 inputs, and the Pareto-selected configuration of Table 2 in the main paper. Statistical tests are paired two-sided Wilcoxon signed-rank tests with Holm correction across metrics within an experiment; effect sizes are paired Cohen’s d .

Contents

S1 Data-augmentation ablation	1
S2 Class-imbalance (pos_weight) sensitivity	1
S3 Selection criterion: Pareto vs. scalar (full results)	2
S4 Architecture-agnostic validation (ResNet-18)	3
S5 Cross-domain generalization (SIDTD)	4
S6 Augmentation does not close the cross-domain gap	5

S1 Data-augmentation ablation

Sec. 4.1 states that standard augmentation degraded every metric on the FantasyID holdout. Table [S1](#) gives the full 30-seed comparison. The augmented model applies, in the training loader only, random rotation ($\pm 10^\circ$), brightness, contrast, Gaussian blur, and JPEG compression, using the *same* Pareto configuration. On the clean synthetic holdout, augmentation significantly degrades every metric (all Holm $p < 10^{-7}$) and inflates variance by 2–5 $\times$, confirming the paper’s no-augmentation choice for this dataset.

S2 Class-imbalance (pos_weight) sensitivity

Sec. 3.1 notes that the mild class imbalance (2.5:1, attack-majority) leaves PR-AUC robust and that a `pos_weight` sweep confirms no significant effect. Table [S2](#) reports the three settings (inverse-frequency 0.397, neutral 1.0, and the attack-favouring 2.518), each over 30 seeds. No metric differs significantly from the neutral baseline after Holm correction (all $p_{\text{Holm}} > 0.5$), so no loss re-weighting is applied in the main model.

Table S1: Augmentation ablation (blind test, 30 seeds). Δ : aug – no_aug; d : paired Cohen’s d ; p_{Holm} : Holm-corrected paired Wilcoxon. All differences significant.

Metric	No aug (ours)	Aug	Δ	d	p_{Holm}
PR-AUC	0.9967 ± 0.0021	0.9577 ± 0.0457	-0.0390	-0.84	1.5×10^{-8}
Dice	0.8750 ± 0.0180	0.7797 ± 0.0396	-0.0953	-2.10	9.3×10^{-9}
B _{Acc}	0.9573 ± 0.0184	0.8302 ± 0.1007	-0.1271	-1.20	5.6×10^{-9}
F1 ₁	0.9651 ± 0.0152	0.8593 ± 0.1061	-0.1058	-0.97	1.1×10^{-8}
F1 _{macro}	0.9423 ± 0.0234	0.7998 ± 0.1136	-0.1425	-1.19	7.5×10^{-9}

Table S2: pos_weight sensitivity (blind test, 30 seeds). Baseline is pos_weight= 1.0. No comparison against the baseline reaches significance after Holm correction.

pos_weight	PR-AUC	Dice	B _{Acc}	Rec. attack	Rec. bonafide
0.397 (inv. freq.)	0.9973 ± 0.0018	0.8793 ± 0.0168	0.9614 ± 0.0146	0.9431	0.9797
1.0 (baseline)	0.9967 ± 0.0021	0.8750 ± 0.0180	0.9573 ± 0.0184	0.9435	0.9712
2.518 (attack-favouring)	0.9967 ± 0.0019	0.8688 ± 0.0240	0.9552 ± 0.0164	0.9435	0.9669

S3 Selection criterion: Pareto vs. scalar (full results)

Table 3 of the main paper reports a compact form of the selection-criterion study. Table S3 gives the complete 30-seed results with 95% confidence intervals for all five metrics, and Table S4 the paired statistics. The two reported objectives quantify the two tasks: PR-AUC (area under the precision–recall curve) is a threshold-independent measure of the detection head, summarizing precision over the full range of recall and remaining informative under class imbalance; Dice (the Sørensen–Dice overlap coefficient, $2|A \cap B|/(|A| + |B|)$ between the predicted and ground-truth forged-region masks) measures localization quality. Our Pareto criterion selects the trial closest to the ideal point (1, 1) in the normalized PR-AUC/Dice plane, whereas **by_prauc** and **by_dice** take the global argmax of validation PR-AUC or validation Dice, respectively.

Significance is assessed with paired two-sided Wilcoxon signed-rank tests over the 30 shared seeds, with Holm correction across the five metrics within each comparison; effect sizes are paired Cohen’s d . Pareto improves localization over **by_prauc** (Dice $\Delta = +0.0161$, $d = 0.59$, $p_{\text{Holm}} = 0.043$) while not degrading detection ($\Delta\text{PR-AUC} = +0.0014$, $p_{\text{Holm}} = 0.062$, n.s.), and improves *both* objectives over **by_dice** (PR-AUC $\Delta = +0.0019$, $d = 0.64$, $p_{\text{Holm}} = 0.0072$; Dice $\Delta = +0.0185$, $d = 1.00$, $p_{\text{Holm}} = 0.00039$), the latter a large effect. All gains over **by_dice** except F1₁ reach significance after correction (Table S4).

Table S3: Selection-criterion comparison (blind test, 30 seeds). Mean $\pm$ Std with 95% CI in brackets.

Metric	Pareto (ours)	by_prauc	by_dice
PR-AUC	0.9967 ± 0.0021 [0.9959, 0.9975]	0.9953 ± 0.0021 [0.9945, 0.9961]	0.9948 ± 0.0027 [0.9938, 0.9958]
Dice	0.8750 ± 0.0180 [0.8683, 0.8818]	0.8589 ± 0.0230 [0.8504, 0.8675]	0.8566 ± 0.0147 [0.8511, 0.8620]
B _{Acc}	0.9573 ± 0.0184 [0.9504, 0.9642]	0.9481 ± 0.0146 [0.9427, 0.9536]	0.9463 ± 0.0170 [0.9400, 0.9527]
F1 ₁	0.9651 ± 0.0152 [0.9594, 0.9708]	0.9569 ± 0.0122 [0.9523, 0.9614]	0.9564 ± 0.0121 [0.9519, 0.9609]
F1 _{macro}	0.9423 ± 0.0234 [0.9335, 0.9510]	0.9294 ± 0.0184 [0.9225, 0.9362]	0.9283 ± 0.0184 [0.9214, 0.9352]

Table S4: Paired statistics, Pareto vs. each scalar criterion (30 seeds). Δ : Pareto – scalar; d : paired Cohen’s d ; p_{Holm} : Holm-corrected paired Wilcoxon; $\checkmark$: significant at $p_{\text{Holm}} < 0.05$.

Comparison	Metric	Δ	d	p_{Holm}	sig.
Pareto vs by_prauc	PR-AUC	+0.0014	0.49	0.062	
	Dice	+0.0161	0.59	0.043	$\checkmark$
	BAcc	+0.0092	0.40	0.061	
	F1 ₁	+0.0082	0.45	0.066	
	F1 _{macro}	+0.0129	0.46	0.050	$\checkmark$
Pareto vs by_dice	PR-AUC	+0.0019	0.64	0.0072	$\checkmark$
	Dice	+0.0185	1.00	0.00039	$\checkmark$
	BAcc	+0.0110	0.45	0.040	$\checkmark$
	F1 ₁	+0.0087	0.44	0.066	
	F1 _{macro}	+0.0140	0.47	0.043	$\checkmark$

S4 Architecture-agnostic validation (ResNet-18)

Sec. 4.6 states that the identical MOTPE protocol, applied to an ImageNet-pretrained ResNet-18 backbone, yields comparable blind-test performance and that neither model dominates. This section gives the full experimental detail.

Protocol (identical to the main paper). We run a single MOTPE study with exactly the search space of Table 1 (learning rate, weight decay, dropout, decoder width C_{dec} , and λ_{mask}), the same 5 inner folds, the same 50 trials, the same 15 epochs per trial, the same MOTPE sampler, and the same Pareto selection by minimum Euclidean distance to the ideal point (1, 1). To bound cost we run one study (the outer-fold-0 analogue) rather than the full 10-fold nested loop. The Pareto-selected configuration is then retrained on the same blind holdout over 30 seeds, with early stopping (patience 12 on the validation distance to the ideal point) and the classification threshold fixed from the development validation split, never from the holdout, exactly as for DocVerify. The study is in-memory and writes no weights; it only reads DocVerify’s released blind-test CSV as the paired baseline, so it leaves every paper result untouched.

Architecture integration. Only the encoder changes: the six-block Patel CNN is replaced by a torchvision `resnet18` initialized with the `ResNet18_Weights.IMAGENET1K_V1` checkpoint, downloaded automatically from <https://download.pytorch.org/models/resnet18-f37072fd.pth>. ImageNet normalization (mean [0.485, 0.456, 0.406], std [0.229, 0.224, 0.225]) is applied inside the forward pass (inputs arrive in [0, 1] as in the main pipeline). The five encoder taps are the stem (112×112 , 64 ch), `layer1` (56×56 , 64 ch), `layer2` (28×28 , 128 ch), `layer3` (14×14 , 256 ch), and `layer4` (7×7 , 512 ch, bottleneck). The 512-d global-average-pooled bottleneck feeds the *same* four-layer classification head ($512 \rightarrow 32 \rightarrow 16 \rightarrow 16 \rightarrow 1$, dropout and LeakyReLU as in the paper). For localization, a 1×1 convolution projects the bottleneck to C_{dec} , and four U-Net decoder blocks ($7 \rightarrow 14 \rightarrow 28 \rightarrow 56 \rightarrow 112$) concatenate the `layer3`, `layer2`, `layer1`, and stem skips in turn; a final bilinear upsampling to 224×224 and a 1×1 convolution produce the mask logit. The decoder and both heads are identical to DocVerify; only the feature extractor differs.

Table S5 lists the complete configuration applied to both studies: the fixed training settings, the MOTPE protocol, the blind-test retraining, and the search space with the Pareto-selected value for each backbone. The two studies share everything except the encoder and the resulting selected values.

Table S5: Complete hyperparameter configuration for the DocVerify and ResNet-18 MOTPE studies. All fixed settings and the optimization protocol are identical; only the encoder and the Pareto-selected values differ. Stratified group k -fold and MOTPE sampler are both seeded at 42.

Fixed training settings (both backbones)		
Input	224 × 224, normalized to [0, 1]	
Batch size	64	
Optimizer	AdamW	
LR scheduler	ReduceLROnPlateau (factor 0.5, patience 3)	
Mixed precision	AMP (FP16/FP32)	
Gradient clip	1.0 (L2 norm)	
Activation	LeakyReLU ($\alpha = 0.2$, fixed)	
Mask threshold	0.5	
Loss	$\text{BCE}_{\text{cls}} + \lambda_{\text{mask}} (\text{BCE} + \text{Dice})_{\text{seg}}$	
MOTPE protocol (both backbones)		
Inner folds	5 (stratified group k -fold)	
Trials	50 (MOTPE sampler)	
Epochs / trial	15 (validate every 2, patience 1)	
Selection	min. distance to ideal point (1, 1)	
Blind-test retraining (both backbones)		
Epochs	up to 100, early stopping patience 12	
Seeds	30 (seeds 42–71)	
Threshold	from dev. validation split (never holdout)	
Search space and Pareto-selected value		
Parameter	Range / set	DocVerify / ResNet-18
η	$[5 \times 10^{-5}, 9 \times 10^{-4}]$ log	5.21×10^{-4} / 3.68×10^{-4}
λ_{wd}	$[10^{-7}, 10^{-4}]$ log	4.27×10^{-5} / 1.45×10^{-6}
p (dropout)	[0.1, 0.4] uniform	0.103 / 0.381
C_{dec}	{96, 128, 192, 256}	192 / 256
λ_{mask}	[0.5, 3.0] uniform	2.46 / 2.69
Distance to ideal	0.316 (DocVerify) / 0.164 (ResNet-18)	

Result. Table S6 gives the full 30-seed comparison. The ResNet-18 Pareto-selected configuration (distance to ideal = 0.164) is $\eta = 3.68 \times 10^{-4}$, $\lambda_{\text{wd}} = 1.45 \times 10^{-6}$, dropout = 0.381, $C_{\text{dec}} = 256$, $\lambda_{\text{mask}} = 2.69$, the last consistent with DocVerify’s 2.46, confirming that the Pareto balance the method discovers is stable across backbones. DocVerify is better on PR-AUC; ResNet-18 is better on Dice and the F1 metrics; balanced accuracy does not differ significantly. Neither model dominates, so the optimization methodology, not a specific encoder, is what generalizes.

S5 Cross-domain generalization (SIDTD)

Sec. 4.6 reports that the FantasyID-trained model is near chance zero-shot on the out-of-domain SIDTD dataset, but recovers strong performance when the same architecture is retrained on the SIDTD training partition. This section details the full adaptation curve.

Target domain. SIDTD (Boned et al., 2024) is a synthetic dataset of ID and travel documents built from real document templates, a distinct visual domain from FantasyID’s invented *fantasy*

Table S6: DocVerify vs. ResNet-18 under the identical MOTPE protocol (blind test, 30 seeds). Δ : DocVerify – ResNet-18; d : paired Cohen’s d ; p_{Holm} : Holm-corrected paired Wilcoxon.

Metric	DocVerify	ResNet18-MOTPE	Δ	d	p_{Holm}
PR-AUC	0.9967 ± 0.0021	0.9942 ± 0.0025	+0.0024	+0.75	2.8×10^{-4}
Dice	0.8750 ± 0.0180	0.8894 ± 0.0111	−0.0144	−0.60	8.1×10^{-3}
BAcc	0.9573 ± 0.0184	0.9629 ± 0.0108	−0.0055	−0.21	0.285 (n.s.)
F1 _l	0.9651 ± 0.0152	0.9779 ± 0.0069	−0.0128	−0.67	3.8×10^{-3}
F1 _{macro}	0.9423 ± 0.0234	0.9613 ± 0.0114	−0.0190	−0.63	4.0×10^{-3}

identity cards. We use only the *templates* subset present on disk (2,222 images: 1,000 bonafide and 1,222 attack), not the full SIDTD video collection. Masks are derived from the provided annotations at the same 224×224 resolution as the main pipeline. The dataset is obtained from its official distribution at <http://datasets.cvc.uab.es/SIDTD/> (doi:10.1038/s41597-024-04160-9).

Zero-shot. The 30 FantasyID blind-test multitask checkpoints are evaluated directly on the SIDTD templates with no adaptation, at a fixed classification threshold of 0.5. We headline PR-AUC and ROC-AUC because they are threshold-independent and therefore measure whether the model still *rank*s bonafide vs. attack across domains; balanced accuracy and F1 at a fixed 0.5 threshold are unreliable cross-domain and are reported only for completeness.

Few-shot and full retraining. For each budget $n_{\text{shots}} \in \{10, 25, 50, 100, 200, 400\}$ we draw a class-stratified sample from the SIDTD training partition and fine-tune the same architecture, repeating over 25 independent sample draws (seeds) per budget to capture sampling variance; the full-retraining point ($n_{\text{shots}} = 1777$) uses an 80/20 stratified split of the templates subset (1,777 train, 445 test). Table S7 gives the resulting curve from zero-shot through few-shot to full retraining. Few-shot adaptation up to 400 samples stays near chance; only full in-domain retraining recovers strong detection (PR-AUC 0.878, ROC-AUC 0.826), confirming that the gap reflects domain shift rather than architectural capacity. Section S6 further shows that training-time augmentation does not substitute for in-domain data.

Table S7: SIDTD adaptation curve. $n_{\text{shots}} = 0$ is zero-shot (30 seeds); $n_{\text{shots}} \geq 10$ are few-shot/full retrainings (25 seeds each). Templates subset (1 000 bonafide, 1 222 attack).

n_{shots}	PR-AUC	ROC-AUC	BAcc	F1 _{macro}
0 (zero-shot)	0.5554 ± 0.0110	0.5041 ± 0.0094	0.5010	—
10	0.5485 ± 0.0084	0.4893 ± 0.0147	0.5002 ± 0.0135	0.4841 ± 0.0199
25	0.5595 ± 0.0089	0.4996 ± 0.0092	0.5046 ± 0.0146	0.4755 ± 0.0255
50	0.5575 ± 0.0102	0.4960 ± 0.0128	0.4953 ± 0.0167	0.4652 ± 0.0207
100	0.5706 ± 0.0208	0.5060 ± 0.0199	0.5032 ± 0.0203	0.4848 ± 0.0259
200	0.5774 ± 0.0170	0.5089 ± 0.0144	0.5031 ± 0.0182	0.4853 ± 0.0289
400	0.5912 ± 0.0238	0.5123 ± 0.0171	0.5028 ± 0.0184	0.4843 ± 0.0349
1 777 (full)	0.8780 ± 0.0161	0.8263 ± 0.0228	0.7188 ± 0.0349	0.7142 ± 0.0333

S6 Augmentation does not close the cross-domain gap

As a follow-up to Sections S1 and S5, we tested whether augmentation-trained checkpoints generalize better zero-shot to SIDTD. Table S8 shows a null result: both the augmented and non-

augmented models remain at chance, with no significant difference in detection quality after Holm correction. Augmentation therefore does not substitute for in-domain data; closing the gap requires multi-source training (future work, Sec. 6).

Table S8: Zero-shot SIDTD generalization, augmented vs. non-augmented FantasyID training (30 seeds each). Δ : $\text{aug} - \text{no_aug}$; p_{Holm} : Holm-corrected paired Wilcoxon.

Metric	No aug	Aug	Δ	p_{Holm}
PR-AUC	0.5554 ± 0.0110	0.5639 ± 0.0096	+0.0085	0.077 (n.s.)
ROC-AUC	0.5041 ± 0.0094	0.5089 ± 0.0060	+0.0048	0.199 (n.s.)

Data availability. The result CSVs underlying every table above are released in the reproducibility archive accompanying this document. The full source code and experiment scripts are at <https://github.com/HernanDiaz/IDVerify>. Datasets and pretrained weights are obtained from their original sources: FantasyID from <https://zenodo.org/records/17063366>; SIDTD from <http://datasets.cvc.uab.es/SIDTD/> (doi:10.1038/s41597-024-04160-9); and the ResNet-18 ImageNet weights from torchvision (`ResNet18_Weights.IMAGENET1K_V1`).